



Institutt for Informasjonsteknologi

Postadresse: Postboks 4 St. Olavs plass, 0130 Oslo

Besøksadresse: Holbergs plass, Oslo

PROSJEKT NR.
23-2022

TILGJENGELIGHET
Offentlig

Telefon: 22 45 32 00

BACHELORPROSJEKT

HOVEDPROSJEKTETS TITTEL IoT - Oppgave	DATO 25.05.2022
	ANTALL SIDER / BILAG Over 9000
PROSJEKTDeltakere Andreas Torres Hansen Uy Quoc Nguyen Anders Hagen Ottersland	INTERN VEILEDER Jianhua Zhang

OPPDRAGSGIVER Accenture	KONTAKTPERSON Daniel Meinecke
----------------------------	----------------------------------

<p>SAMMENDRAG</p> <p>Det er utviklet en proof-of-concept løsning for et sentralisert system der en server hvor selvkjørende biler kan kommunisere med hverandre øker trafikkflyt i et kryss. En semi fysisk demonstrasjon er laget for å vise fram løsningen.</p>

3 STIKKORD Internet of Things
Raspberry Pi
Sentralisert System

DATA3900 - Bachelor project

Final report

Gruppe 23:

Hansen, Andreas Torres (s338851)

Nguyen, Uy Quoc (s341864)

Ottersland, Anders Hagen (s341883)

Total pages: [52](#)

Last updated:

May 24, 2022

Contents

Abstract	1
Preface	2
I Preliminary	3
1 Introduction	4
1.1 Stakeholders	4
1.1.1 Students	4
1.1.2 Product owner	4
1.1.3 Supervisors	4
1.1.4 Client	4
1.2 Project Description	5
1.2.1 Project background	5
1.2.2 Significance	5
1.2.3 Goals and requirements	5
1.2.4 Problem statement	5
II Process Documentation	6
2 Work methodology and technology used	7
2.1 Development method	7
2.1.1 Scrum and sprints	7
2.1.2 Kanban as our backlog	9
2.2 Tools and technologies	10
2.3 Prioritization method	10
3 Phases	12
3.1 Phase 1 - Planning and research phase	12
3.1.1 Choice of programming languages	13
3.1.2 Internet of Things	13
3.1.3 Preventing traffic congestions for a one-lane road	13
3.2 Phase 2 - REST API and why it did not work with our project	16
3.2.1 Alternative solution 1 - Short-polling and long-polling	17
3.2.2 Alternative solution 2 - Webhooks	17

3.2.3	Alternative solution 3 - WebSockets	17
3.3	Phase 3 - WebSocket with SignalR	17
3.3.1	Implementation of traffic on a single lane	18
3.3.2	Extending the concept to a more complex scenario - traffic on an intersection	19
3.4	Phase 4 - Semi physical simulation	19
3.4.1	Building a new car	19
3.4.2	Calibration of the cars	20
3.4.3	Construction of semi physical demonstration	22
4	Conclusion	25
4.1	Results	25
4.1.1	Real world scenario	25
4.1.2	Possible improvements	26
4.2	Further discussions	26
4.2.1	Self evaluation	26
4.2.2	Educational Value	27
4.2.3	Real world application	27
4.2.4	Edge computing and AI	28
4.2.5	System security	28
III	Product Document	29
5	Program description	30
5.1	Essential code snippets behind the system	32
5.1.1	Initialization	32
5.1.2	Handshake and listener	34
5.1.3	Patch	37
5.1.4	VehiclesHubDatabase	38
5.1.5	RoutePlanner and SetTravelPlan	42
A	User manual	45
	Bibliography	45
B	Glossaries and Acronyms	49
B.1	Glossaries	49
B.2	Acronyms	50
C	Sprint documents	51

ABSTRACT

In 2020 a group of students from Høyskolen Kristiania developed self-driving cars with Raspberry Pi for Accenture. This year, our group was tasked to extend their project by combining it with a centralized communication system. We consider how traffic flow can be improved with autonomous cars and a server managing the traffic. Consequently, our group developed a client-server program where the existing Raspberry Pi vehicles communicate with a server to improve traffic flow at an intersection.

Traffic lights control intersections today to prevent accidents. However, stopping and accelerating vehicles lead to traffic congestion (**traffic_shockwave**). Hence, this project concludes, through demonstrations, shows that a centralized server can indeed improve traffic flow by adjusting the vehicles' speed accordingly.

PREFACE

This is the report of our bachelor thesis at Oslo Metropolitan University, Faculty of Technology, Art and Design. We tried to develop a solution for traffic management with self-driving cars and server communication. Our project was created for Accenture and lasted from January 2022 to May 2022. In this project, we documented the functionality and process of making it. We physically demonstrated our proposed solution with a Raspberry Pi computer working as a car.

This report is split into three parts: First, an introductory part where the project is introduced. Then a process documentation, where we elaborate on the process of the development. Lastly, the product documentation elaborates on the structure of the program.

We would like to thank everyone that has contributed to our project. We would especially like to thank:

- Ivar Austin Fauske, Solfrid Johansen and Benjamin Vallestad, representatives from Accenture.
- Dr. Jianhua Zhang, internal supervisor at OsloMet.

PART I

PRELIMINARY

This part serves as an introduction to the project. Throughout this part of the report, we will introduce this project's contributing members and clients. Additionally, we will elaborate on the goals and requirements surrounding our project and introduce the research and development question this project explores.

Chapter 1

Introduction

1.1 Stakeholders

1.1.1 Students

Andreas Torres Hansen, Software Engineering

Anders Hagen Ottersland, Software Engineering

Uy Quoc Nguyen, Software Engineering

1.1.2 Product owner

Benjamin Vallestad

1.1.3 Supervisors

Professor Jianhua Zhang, Internal Supervisor

Ivar Austin Fauske, External Supervisor

Solfrid Hagen Johansen, External Supervisor

1.1.4 Client

Accenture AS

Accenture is an international IT firm that operates in 200 different countries worldwide. They offer a wide range of services in many fields, like artificial intelligence, data analytics, and cloud computing. The main office is in Dublin, and the Norwegian main office is in Fornebu. Accenture has 674 thousand employees internationally, of which 1000 work in Norway (**accenture_earning_report_2021**). In 2021, Accenture generated a revenue of approximately \$50.3 billion (**accenture_about**).

1.2 Project Description

1.2.1 Project background

For many years, Accenture has offered innovative bachelor thesis projects for students at OsloMet and Høyskolen Kristiania. In 2020, a group of students from Høyskolen Kristiania was developing model-sized self-driving vehicles using Raspberry Pi in conjunction with machine learning as their project. This year our group was offered to extend this project further; to explore plausible improvement with the addition of a centralized communication system.

A centralized communication system follows a client/server architecture, where multiple clients connect to a server, and the server handles all the primary processing of data (**centralized**). In this project, the Raspberry Pi vehicles act as a client that connects to a server that performs the major decisions for the vehicles. Thus, this report uses the terms; vehicles, and clients interchangeably, while the server refers to the centralized server.

According to our product owner, Norway is one of the countries that are ready to utilize self-driving cars. However, self-driving vehicles alone are likely not enough to solve all of today's traffic challenges. Hence, this project aims to solve the issue by introducing a management system for autonomous vehicles and evaluating the value such a system can provide.

1.2.2 Significance

Our product owner states that this project provides value for Accenture in the shape of building knowledge around new technology and theory. Accenture also wants to explore the potential positive societal and climate effects a centralized communication system for transportation could provide.

1.2.3 Goals and requirements

The project aims to produce a prototype that can be shown to stakeholders and, either physically or digitally, demonstrate how Accenture could combine self-driving cars with a centralized system. We will utilize the self-driving vehicle built for Accenture from the previous bachelor thesis. Preferably, the prototype can show at least one situation where the outcome differs depending on self-driving cars plus a centralized system versus only using self-driving vehicles. Furthermore, the cars should also be able to drive independently with the incorporated AI model if the centralized system is out of reach. The system should also be scalable so that more vehicles can be added or removed at a later point in time.

1.2.4 Problem statement

Based on the goals and requirements set by Accenture, we have formulated the research and development question:

“How can we improve traffic flow by using a combination of self-driving cars and a centralized communication system?”

PART II

PROCESS DOCUMENTATION

The process documentation involves the process of the project, and elaborates on the challenges and solutions that the group has explored. Moreover, it details how the group has worked and what the group did to make this project a success.

Chapter 2

Work methodology and technology used

2.1 Development method

We had very few requirements and technical restrictions when we received the project, which left the project open to interpretation. Therefore, we wanted to choose a flexible work methodology. Agile work methods focus on continuous planning throughout the process and having frequent communication with the client, in our case Accenture. We had meetings with our external supervisors from Accenture once every second week, which meant the agile model was a good fit for our project.

We took inspiration from two light frameworks, Kanban and Scrum. Scrum is an agile, light framework that helps people and teams work together. Scrum describes a set of meetings, roles, and tools.

A sprint is an essential part of using the Scrum framework. Sprints are a fixed time length, often between one and four weeks. In this specified time length, the teams do tasks assigned from the sprint backlog. Each sprint starts with sprint planning and ends with a sprint retrospective. We found it most viable for our project to plan in increments of two weeks. We chose two-week increments because it gave us a good balance between work and planning. We also had meetings with our client Accenture every two weeks, which fitted well with the time increment.

2.1.1 Scrum and sprints

Scrum is a framework that dictates how developers work in teams to solve complex problems. The development process is also divided into time intervals called a sprint. A sprint is an essential part of using the Scrum framework (**prosjektveilederen**). Sprints are a fixed time length, often between one and four weeks. In this specified time length, the teams do tasks assigned from the sprint backlog. Each sprint starts with sprint planning and ends with a sprint retrospective.

Our group also used the meetings in the Scrum framework, which consist of sprint planning, sprint retrospective, and daily standups. Sprint planning is a meeting or event which starts before a sprint. During sprint planning, teams agree on goals for the sprint and what tasks from the backlog should be prioritized. The backlog is a list of functionality the product should contain. In addition, we wrote down the tasks for the specific sprints in Google Docs. These tasks were to be finished by the next sprint. [Figure 2.1](#) shows our sprint planning document that lists our goals for each sprint.

Sprints for bachelorproject	
<u>Sprint id</u>	<u>Oppgaver</u>
1	Implementere web API Lagring av informasjon til database Organisering av prosjektet på git-hub
2	Integrere web-api løsningen vår med bilene fra forrige prosjekt Implementere enhetstesting Gjøre mer research for løsninger på oppgaven(webhooks, kafka, signalR)
3	Gjøre mer research på webhooks Få bilene til å sende riktig informasjon til server Implementere signalR server
4	Starte å skrive på bacheloroppgaven Få server til å sende kommandoer til flere biler Få bilene til å reagere på kommandoene til serveren
5	Bygge bil nr 2 Lage et veisystem for demo med veikryss Skrive et førsteutkast for rapporten
6	Lage en demo som viser IoT-systemet Bli ferdig med å bygge bil 2 Skrive mer om implementasjon og starte å skrive på løsningen på rapporten
7	Gjøre ferdig alt det tekniske ved oppgaven Ferdigstille demoen Skrive ferdig implementasjon og løsnings-delen på oppgaven

Figure 2.1: An overview of our sprints. Each sprint lasted 2 weeks. This figure shows the main tasks of each sprint.

After a sprint, we would have a sprint retrospective to discuss what went well and what we could have done better. We would also examine if the task assigned in the sprint meetings was finished or needed more work. These meetings helped us reflect over the prior week and adjust accordingly, if necessary. The sessions also helped us determine if we were on track with our initial plan.

In addition to the weekly meetings, we also had daily standups. Daily standups are short meetings, usually lasting around five minutes, where each person answers three questions:

- What did the person do last time?
- What is the person going to do today?
- Are there any challenges?

We implemented daily standups because it helped our team get on the same page, and it made it easier to plan what each of us had to do that specific

day.

Scrum often consists of a team with different roles. As a team of three, we did not feel the necessity to have specified roles because we usually worked together on our projects. However, we alternated on being the scrum master. The scrum master's responsibility is to keep track of the backlog and lead the sprint planning meetings.

2.1.2 Kanban as our backlog

Our implementation of Kanban was to use a Kanban board as the backlog. We used a Kanban board to visualize where a task is in the work process. Figure 2.2 shows an example from our project, with description labels and priority labels:




Figure 2.2: Extract of our Kanban board from Github. Each task has a color-coded label representing the priority, and a label to describe the task.

We have four columns that represent which phase a task is in. The backlog is the tasks in the to-do column. We dragged it over to the "In progress"-column when we worked on a task. After finishing a task, it went to the "Review in progress"-column, where other group members reviewed the code. If we concluded that the task was completed in the reviewing, it was moved to the "Done"-column. When creating a task, we tag it with a high, medium, or low to indicate its priority. In addition, to the priority tag, we gave each task a label; feature, bug, and non-technical to communicate what the task entails. One feature of the GitHub project is to connect tasks to specific branches, so the task automatically gets finished when merging the branch into the main branch. The Kanban board was a great tool to see which tasks to choose for our sprints and also keep track of where the tasks were in their process.

However, we did not use the Kanban board throughout the whole process. The backlog changed a lot, and the Kanban board needed many modifications to

be up to date. We figured out that it was more beneficial to focus on one framework, which in our case was Scrum. In addition, we were able to keep track of the tasks by having frequent meetings.

2.2 Tools and technologies

The circumstances surrounding Covid 19 meant that we could not meet our supervisors in person at the start. Luckily, the group could still meet physically a few times a week. We had to use a wide range of tools for communication. Most restrictions were removed later in the project, but we kept our meetings with our supervisors digital throughout the project.

- Email - Formal communication with supervisors and product owner
- Teams - Meetings and the platform of choice for communicating with the external supervisors on an informal level.
- Zoom - Meetings with our internal supervisor at OsloMet
- Facebook Messenger - Communication internally in the group. We used it to send messages to each other when we were not physically together and to send pictures of code.

The project required us to collaborate while working on different computers, leading to potential overlapping. Therefore, tools that helped us work on the project together were essential. We used the following tools for this purpose:

- Git and Github - Version control of choice
- Google docs - Used to write our journal and other documents that need to get updated regularly, and to share documents.
- Latex - Used to write reports.

We also needed text editors that supported the programming languages that we used. Our solution was created with Python and C#.

- Pycharm - IDE for coding in Python
- Visual studio and Rider - IDE for coding in C#
- Thonny - Text editor for coding in Python on Raspberry pi

Project planning and documentation were also an important part of the project. The tools we used for the project planning were:

- Github project - Kanban board and creating backlog tasks
- Excel - Used for visualizing our work plan with tables and a Gantt diagram

2.3 Prioritization method

The MoSCoW method is a prioritization technique used in project management. The word MoSCoW is an acronym where:

- "Mo" stands for must-have and represents our project's most prioritized requirements. These are necessary for the success of our project.
- "S" stands for should have. Our project should include these features, but they are not mandatory.
- "Co" stands for could have. We want to include these features in our project, but they are not prioritized.
- "W" stands for will not have this time. The features in the "W" section might be for a later group if someone wants to build on our project further.

Since we were unsure of how many features we could finish within the time frame of the project period, we thought a prioritization method was a good fit. [Figure 2.3](#) shows our visualization of the Moscow method.



Figure 2.3: Visualization MoSCoW method. Everything under "Mo" are requirements we must have in our project. Under the "S" sections are features we should have. Under "Co" are features we should have but are not necessary. Under "W" are features we will not have this time

Chapter 3

Phases

Our sprints each lasted two weeks, as per described in section [2.1.1 Scrum and sprints](#), but in retrospect, it is apparent that we can divide our process into four phases:

Phase 1	- Planning and research phase	week (1-4)
Phase 2	- Rest API and why it did not work with our project	week (5-8)
Phase 3	- Websocket with SignalR	week (9-12)
Phase 4	- Demonstration of our system	week (13-18)

An overview of our sprints are in figure [Figure 2.1](#)

3.1 Phase 1 - Planning and research phase

After we had gotten in touch with Accenture and spoken with the supervisors and the product owner, the group had to make a few decisions regarding the project's direction.

An important choice for the project was to either build and train an AI model for the vehicles from scratch or use the existing model. Building a new AI model would provide a deeper understanding of the model we could utilize. Due to the time constraint, continuing with the existing model was more favorable. In addition, we believed that the project would have the potential to be too similar to the previous project. Our group also had more prior experience with networking than with Raspberry Pi and AI. We, therefore, chose to use the AI model from the previous project.

In this phase, we did not have access to the vehicle, nor the code, made by the prior bachelor group. However, there was a need for project planning and research before we could start developing our IoT system. The topics that needed to be researched were:

- What causes traffic jams and solutions to fix it.
- IoT-systems and how they function with vehicles.

- Planning and development methods that will fit our project

We also used the pre-project phase to get to know Accenture, their guidelines, and their workspace. Moreover, we participated in a "Design thinking"-workshop and made a work contract.

3.1.1 Choice of programming languages

We chose to implement the client in Python and the server in C#. The group before us had used Python for their Raspberry Pi vehicle, making Python a natural choice to extend the code from their project. Our group also had experience with networking in Python. Furthermore, the .NET ecosystem has well-developed solutions for creating IoT applications, microservices, and web applications (**dotnet**). We had to write our server in C#, to take advantage of these solutions. The server needed to be as efficient as possible, and C# is also considered a fast programming language (**csharp**). We also had some prior knowledge of coding in C#.

3.1.2 Internet of Things

The Internet of Things refers to physical objects that communicate using sensors, cameras, software, or other technologies connecting and exchanging data. This communication takes place over the internet or other communication forms. The number of connected IoT devices in the world is increasing, and it is becoming a big part of society (**iot_analytics**). IoT has also been evolving in recent years due to other technologies becoming more accessible, such as machine learning and the 5G network.

IoT projects can, for instance, be applied in climate surveillance systems, energy, or transportation. In this thesis, we will explore the possibilities of using IoT in transportation, more specifically in personal automobiles. The convergence of these fields is more commonly known as IoV, Internet of Vehicles. An IoV system is a distributed system for wireless communication and information exchange between vehicles through agreed-upon communication protocols (**chinese_iov**). The system could potentially integrate functionality for dynamic information exchange, vehicle control, and smart traffic management. In our thesis, we will explore these possibilities on a small scale.

3.1.3 Preventing traffic congestions for a one-lane road

Traffic congestion, also known as traffic jams, is when a long line of vehicles moves slowly or has stopped moving altogether. Traffic jams can create frustration and disrupt nearby local environments with sound and gas emissions (**traffic_congestion_pollution**). Many factors can cause traffic congestion, such as: poorly designed roads, not wide enough roads, traffic light patterns, and accidents (**traffic_congestion**).

With this in mind, we started by focusing on a simple scenario: when a car drastically reduces its speed or completely stops on a single-lane road.

This scenario will lead to the vehicles behind needing to slow down drastically as well. This phenomenon is called traffic jam shockwave (**traffic_shockwave**).

To prevent this, we propose a solution where cars reduce their velocity before they reach the destination of where the shockwave started. For this to happen, a server could keep track of the cars' positions and send information to the vehicles behind, when required.

We came up with an idea on how the server and cars should interact. First, the car would connect to the server and provide information about its current speed, weight, width, and length. The server will use this information to keep track of all the cars' positions on the road. The cars would send information to the server if their velocity changed. This message would trigger an event on the server where it would command all the cars behind to slow down accordingly. [Figure 3.1](#) shows a flow chart of a potential simulation of this solution:



Figure 3.1: This figure shows the flow diagram of our first proposed solution.

3.2 Phase 2 - REST API and why it did not work with our project

The project aimed to connect Raspberry Pi devices to a server. Our initial approach was to implement a RESTful API on a server, as the connection layer between them. We also considered data security a critical aspect of an upscaled version of such a system. The group discussed how to store the data about the vehicle's positions during this phase, keeping in mind that we wanted our proof-of-concept to be scalable to real life.

Representational state transfer (REST) *application programming interface* (API) provides a way for clients and servers to establish communication through *hypertext transfer protocol* (HTTP), which is a protocol for transferring data between network devices. Using a REST API and the standard HTTP methods; POST, GET, PATCH, DELETE, clients can send requests to a server to perform standard CRUD (create, read, update and delete) operations on a database respectively (**rest_api**). We wanted to keep track of all connected cars, on a database on the server. Therefore, REST API seemed like a good fit for our current solution.

Due to the time constraint of the program development, quickly choosing an appropriate type of database for our system was desirable, as migrating databases later could be a big timesink. We chose to incorporate a time-series database, which automatically includes a timestamp for each database entry. InfluxDB is a time-series database created by InfluxData and provides SQL-like syntax that is quick to query resources. Moreover, it provides both ease of installation and supports a multitude of languages (**influxdb**).

During this phase, the group implemented a standard REST API server with C# with the idea that the Raspberry Pi vehicles should exchange information on their velocities, accelerations, and positions to the server. When initially connecting to the server, the client should first provide the information of its properties, through a POST call. The server will then add the vehicle to the InfluxDB to keep track of the vehicle's information. The vehicle should be able to perform a GET request to the server to retrieve its information.

At this stage, the client will send PATCH requests to the server to update its information. In addition, the server will add a new entry to the database whenever it receives this request. The idea was that the client and server would be able to continuously communicate with each other, creating a live feed of the clients' behavior.

However, it became apparent that a RESTful API could not achieve the results that we wanted. Firstly, the server was only able to communicate with one client at a time, i.e., the client that sends a request. The group wanted the server to respond to other clients without a request. In our case, whenever a vehicle sends a PATCH request, the server should be able to inform other vehicles of this event.

A new solution had to be in place to achieve our goal, so we consulted our supervisors for help. They proposed some solutions for us to research, including long-polling, Webhooks, and Websockets discussed in [3.2.1 Alternative solution](#)

[1 - Short-polling and long-polling](#), [3.2.2 Alternative solution 2 - Webhooks](#) and [3.2.3 Alternative solution 3 - WebSockets](#), respectively.

3.2.1 Alternative solution 1 - Short-polling and long-polling

Polling refers to the server pushing resources to the client. There are mainly two types of polling; short- and long-polling.

When short-polling, a client requests a resource from the server, and the server responds with an empty response if the resource is not available. The client will then send a new request after a short amount of time, and the cycle repeats until the client receives the resource it has requested.

Long-polling is similar to short-polling, but the server does not send anything back until the resource is available. In other words, the client sends a request to the server, and the server holds this request until it has a response available to the client. In our case, we wanted every client to perform a GET request to the server. Then the server holds onto this request until it has further instructions for the requesting client.

3.2.2 Alternative solution 2 - Webhooks

Webhooks, according to **webhooks**, is a user-defined callback over HTTP. In our case, implementing webhooks to post notifications on clients based on events sent to the server. This was a good contender to solve our issue. However, implementing webhooks includes extensive research into a system the group had never heard of, in addition to scarce information on how to create such a system. The group decided that the time constraint of this project did not justify the time it would take to implement such a system. We found little information on how to utilize webhooks in our project.

3.2.3 Alternative solution 3 - WebSockets

Websockets is a protocol that provides a bidirectional communication between clients and server by establishing a single TCP connection in both direction (**rfc_websockets**).

3.3 Phase 3 - WebSocket with SignalR

After extensive research on how to solve the two-way communication discussed in [3.2 Phase 2 - REST API and why it did not work with our project](#), the group agreed that WebSockets would be an excellent solution to our problem. Using WebSockets, both the client and server can transfer data at any time. However, **microsoft_websockets** discourages developers from implementing raw WebSockets for most applications and instead recommends using it with SignalR.

ASP.NET SignalR is a library that provides real-time communication using WebSockets while also providing other transport methods such as long-polling, as fallback (**microsoft_signalr**). Furthermore, SignalR API supports *remote*

procedure calls (RPC) using hubs, meaning we can invoke program instructions on the client from the server and vice versa (**microsoft_signalr**).

We decided to implement SignalR on our server, instead of Rest API and InfluxDb, for the reasons given in [3.2 Phase 2 - REST API and why it did not work with our project](#). We used SignalR to implement an echo server, which returns, or "echos", all messages it receives back to the sender. Simultaneously we worked on implementing the client code, so that it could communicate with the server. The client code had to be implemented independently of the Raspberry Pi code. Because our goal was to create a communication module that is reusable through inheritance for other devices, e.g., traffic lights, should it be required to set up a new hub with other devices.

After successfully implementing all the necessary methods on the client, the vehicle class representing the Raspberry Pi device became our next priority. By inheriting the client class, the vehicle class also inherits its ability to connect, listen and send data to the server. Furthermore, the client can also subscribe to events that the server can trigger using RPC.

3.3.1 Implementation of traffic on a single lane

After witnessing a successful connection between the Raspberry Pi vehicle and our SignalR server, we started to implement the necessary functionality on the server, for the simple scenario described in [3.1.3 Preventing traffic congestions for a one-lane road](#). The client will inform the server whenever its velocity has changed. Then, the server will relay this change to every other vehicle behind it on the same road and adjust their velocities accordingly. This server functionality also depends on the continuous monitoring of each vehicle's position. Thus, raising a new issue on how the vehicle information should be stored.

InfluxDB could, in theory, be used to store the vehicle's position; however, since the position is constantly changing, it would require the server to read and write on the database continuously. Hence, the group concluded that this would potentially impact latency on the server. We determined that a live database stored in memory, with vehicles in lists, would be preferable. With simple mathematics, the server could recalculate the vehicle's position based on its previous velocity and a stopwatch whenever it retrieves the information about a vehicle instead.

We then determined to show the product owner from Accenture what we had been working on, thus inviting him and the external supervisor to the next meeting. This short demonstration consisted of running multiple clients with the server to simulate how several vehicles would behave with the regulation of our server. The simulation successfully demonstrated that vehicles were able to adjust their velocities in order to prevent the shockwave phenomena described in [3.1.3 Preventing traffic congestions for a one-lane road](#), and we received positive feedback. However, our supervisors challenged us to extend our concept and make our solution more complex. They believed implementing an intersection where multiple cars meet could show more advanced traffic management.

3.3.2 Extending the concept to a more complex scenario - traffic on an intersection

Expanding further on the concept of regulating speed on a single-lane road, new functionalities on the SignalR server were developed to handle vehicles approaching an intersection. With a more complex topology, expanding the database to account for the new road network was also required. Consequently, code for roads and intersections became needed additions to our solution. The vehicle model on the serverside now also required a route planner to represent what it means to be approaching an intersection.

With these improvements and new functionalities, the server could now command the clients to adjust their velocity to avoid collisions between vehicles approaching an intersection simultaneously. By reducing the velocity of some vehicles, it also became apparent that traffic flow was improved, in contrast to stopping a vehicle.

3.4 Phase 4 - Semi physical simulation

At this point in the development, we were sure about what kind of situation we wanted to simulate to make a satisfying product for Accenture and answer the problem statement in [1.2.4 Problem statement](#). Following the work requirements, we still needed to make a demonstration that showcases how the system works. Starting this work phase, the group and our internal supervisor strongly considered demonstrating our solution virtually. A virtual simulation could produce more data and consequently build a stronger case to answer our problem statement in [1.2.4 Problem statement](#). Contrary, our supervisors at Accenture convinced us that a semi physical simulation utilizing the existing Raspberry Pi vehicle was more in line with Accenture's goals for the project. A semi physical simulation is, according to [Chen2019](#), a simulation that "connects the virtual and real networks, providing an effective means for network design and analysis". For us, this means combining the physical car with a digital server. Hence, the group started to build a new Raspberry Pi vehicle for this purpose.

3.4.1 Building a new car

The previous group had only built one car for their project. Consequently, we needed to build a new car with the remaining components left from the previous project to show a situation where two cars meet at an intersection. However, we only had one TPU, the Coral Usb Accelerator. The TPU is an essential component for giving extra processing power to the computer, and it was necessary to run the artificial intelligence the previous group had used ([prev_project](#)).

Without this accelerator, we could not run the artificial intelligence the previous group had incorporated into their solution. Due to the global chip shortage caused by the Covid pandemic, the accelerator was unavailable to purchase anywhere. As a result, the new car's camera and distance measuring sensor were absent. Not having two cars that utilized artificial intelligence could be a challenge. One of the required features of our solution was that the cars should

be able to drive using the AI when a server connection was unavailable, as per described in [1.2.3 Goals and requirements](#). We decided that as long as we had one car that could navigate traffic independent of the server, the other car could drive solely on commands given by the server. With the components and the product documentation of the previous project [1.2.3 Goals and requirements](#), we were able to build a copy of the car as seen in [Figure 3.2](#).




Figure 3.2: Picture of the two cars we built. To the left is the car without a camera. To the right is the car with a camera on top and a Coral USB Accelerator Edge TPU. This car can take advantage of the onboard AI.

3.4.2 Calibration of the cars

After building a second Raspberry Pi vehicle, we did some testing to figure out how the vehicles behaved together with the server when approaching an intersection. In this test, the vehicles were given a specific velocity and distance by the server. When the vehicles arrived at their destination, the server would tell them to stop and terminate their connection.

The vehicles were able to send information and respond correctly to the server's commands. We also observed that the vehicles drove a different length for each velocity given even though the length given by the server was the same. We assumed that the power given to the vehicle's motors equated to the velocity of the vehicle, which resulted in the unexpected behavior previously mentioned. The demonstration depended on the car's ability to input the correct amount of power to drive at specific speeds instructed by the server. Thus, the group started to measure the distance and time the cars would drive, given intervals of power to determine the relationship between power and velocity.

The measurement setup is as follows: The server instructs the cars to drive with

a given power, which is what the server believed to be the velocity. After a time, the server will stop the car. The group then measured the distance the car had traveled. Given the measured distance and the time given by the server, the real velocity was obtained by dividing distance by time. [Table 3.1](#) shows all the recorded measurements with this set up.

Power	Length (cm)	Time (s)	Velocity (cm/s)
40	467	8.98	52.00
50	425	7.28	58.38
60	400	6.06	66.01
70	357	5.18	68.92
80	325	4.49	72.32
90	314	4.03	77.92
100	286	3.62	79.01

Table 3.1: Measurements taken to determine the relationship between power and velocity. Units were provided to the different parameters, while the power unit was unknown.

We then made a graph in Excel by plotting our data with velocity on the y -axis and power on the x -axis as shown in [Figure 3.3](#).




Figure 3.3: Graph of velocity as a function of power

From [Figure 3.3](#), it appears that the power and velocity were linearly correlated. Hence, using linear regression, we obtained an approximate relationship between power and velocity. [Figure 3.4](#) shows the result of the linear regression performed by Excel using the same dataset in [Table 3.1](#).

Field: **Power** and Field: **Velocity**
appear highly correlated.

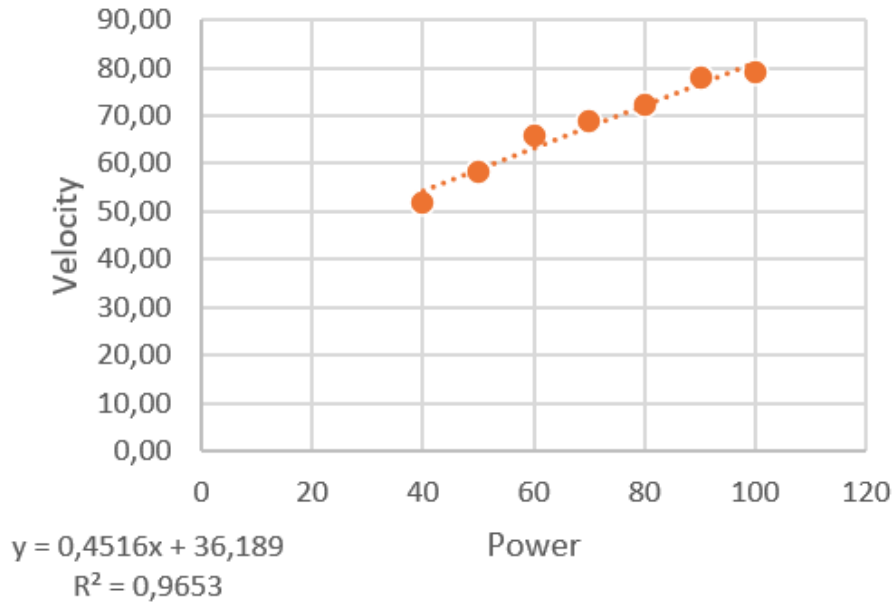


Figure 3.4: Graph of velocity as a function of power with linear regression in Excel. The x-axis shows the Power and the y-axis shows the velocity

From [Figure 3.4](#), the relationship between velocity and power is as follows:

$$v(P) = 0.4516 \cdot P + 36.189 \quad (3.1)$$

where P is power and v is velocity, with a goodness-of-fit measure of $R^2 = 0.9653$. We inserted the relationship into the vehicle and performed a new test to confirm the relationship. We observed that the vehicles drove more or less the correct distance for the desired velocity. Although the distance has some variance in the new tests, we concluded it was accurate enough for our demonstration. If, however, a more accurate formula is desired, then the test can be performed with smaller power intervals.

3.4.3 Construction of semi physical demonstration

We want to test that a centralized communication system and artificial intelligence can improve traffic flow. In order to test the solution we have worked on, we made a physical demonstration where two cars approach an intersection simultaneously. We wanted to observe if the velocity of the vehicles were not drastically changed and, therefore, decreased shockwave described in [3.1.3 Preventing traffic congestions for a one-lane road](#).

We found a space at Accenture that was big enough to build the track. The power input on the Raspberry Pi vehicles is limited to between 40 and 100 units. Since the server will adjust the vehicle's speed to avoid a collision, the intersection was required to be at least 130cm away from the starting point to prevent the server from adjusting speed outside the equivalent power limit. Furthermore, the server prevents cars from driving before at least two vehicles have established a connection. Hence, both vehicles will start their journey simultaneously. We also placed the vehicles at the same distance from the intersection on their respective roads. Given these initial conditions, both vehicles are supposed to collide without the server's intervention. We also made the server log the velocity sent to the cars to track how much the cars' velocities changed.

Making the demo one hundred percent accurate was not possible in our circumstances due to the limitations of the Raspberry Pi. Furthermore, the vehicles were not able to receive the messages simultaneously. Consequently, this meant that they could start with a minor time difference. Another factor was that the trajectory of the vehicles was not always straight. However, we were able to get a consistent demo with enough margins. [Figure 3.5](#) shows an example of a collision during a test demonstration.




Figure 3.5: Here, we can observe the two cars colliding in the intersection. One of the cars started about 5 centimeters further behind the other car and started a few milliseconds later, resulting in a collision. Meanwhile, the server assumed they started simultaneously at the same distance from the intersection.

Furthermore, the server did not account for the lengths of the vehicles during its calculations. After we introduced the length of the vehicles and buffer zone, we were able to get a consistent semi-physical demonstration.

When both vehicles had connected to the server, the cars would drive with an initial speed of 80 cm/s, the upper limit of the Raspberry Pi. Not long after they started to drive, the server recognized the cars approaching the intersection. The server calculates using the car's velocity, position, and length. Then the server calculates which car has to slow down and how much the car needs to slow down to avoid a collision, in this case, to 55 cm/s. After the other car has supposedly passed the intersection, the car that slowed down gets told by the server to speed its velocity back to 80 cm/s. [Figure 3.6](#) is a snapshot of a successful test demo.

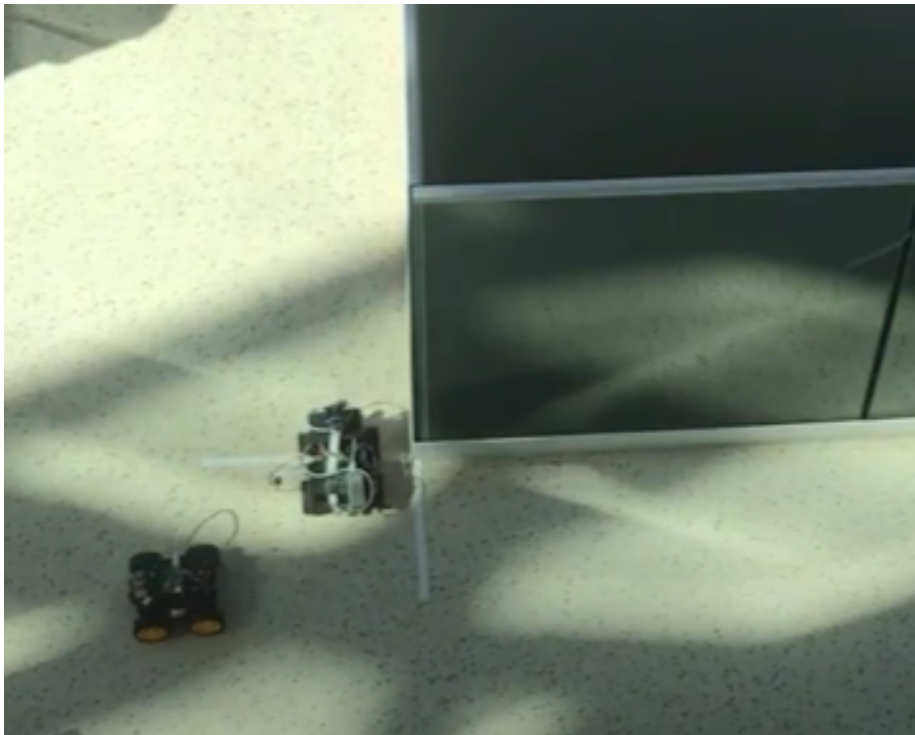


Figure 3.6: Here is a snippet from a successful demo. The car to the left has just passed the intersection, marked as a square with white tape. The car furthest up is, therefore, about to adjust back to its original velocity. Here we can observe that the server prevents a collision.

Chapter 4

Conclusion

4.1 Results

In regards to the problem statement in [1.2.4 Problem statement](#) and the goals for this project discussed in [1.2.3 Goals and requirements](#), we have been able to successfully create a centralized system that communicates that can communicate with autonomous vehicles. A significant change in velocity can lead to a disruption of traffic flow, as discussed in [3.1.3 Preventing traffic congestions for a one-lane road](#). Through multiple semi-physical demonstrations discussed in [3.4.3 Construction of semi physical demonstration](#), it was apparent that the server was able to contribute to an improved traffic flow by preemptively reducing vehicles' to the necessary speed to prevent collisions. Even though this project did only account for scenarios involving intersections, it is also from the exploration done throughout this project that autonomous vehicles, together with a centralized system, can be the answer for improving traffic flow in the future.

Furthermore, we received positive feedback after showing our supervisors and product-owner video snippets of the conducted demonstrations. Additionally, the product owner confirmed that our results adhered to the requirements set by Accenture.

4.1.1 Real world scenario

By comparing the semi-physical demonstration to a real-world scenario where instead of 80 cm/s, the cars would drive 80 km/h before an intersection. In the real world, one of the cars would have to stop before a traffic light while the other could drive past the intersection. Hence, it will involve some cars moving while other stops. In our scenario, the car would only have to slow down to 55 km/h.

Furthermore, our project does not consider other factors, such as curved roads, human mistakes, and animals jumping onto the road. Hence, our demonstrations, conducted in a controlled environment, are not a good reflection of a real-world scenario.

4.1.2 Possible improvements

All the requirements in the MoSCoW method from "could have" to "won't have this time" are requirements for a future project, either for Accenture to improve or for a future bachelor project. Although it fulfilled the product requirements given by Accenture, there is room for improvement. In addition, there is room for improvement with the accuracy of the demo, specifically by doing more calibration tests and making a more accurate formula than we were able to produce. Furthermore, changing the wheels and installing an edge TPU to exploit the existing AI model will result in a more predictable trajectory of the vehicles.

Our demo only contains two cars, but our server supports scenarios with multiple cars. There are also possibilities to connect other devices to the server, for example, traffic lights. However, there are no specific functionalities regarding traffic lights on the server. Scaling the IoT system for functionalities with traffic lights is a task for future development. Adding roads and making a more complex road system would also be a task for future development.

As mentioned, we made a semi-physical demonstration. However, a virtual simulation will yield more data for research while also exploring other more complex road configurations.

Making a viable product in the real world is a long way ahead. That will require a lot more testing and implementation on a bigger scale. However, we hope that our testing and research can be of value towards that step.

4.2 Further discussions

In addition to implementing the program, we also had discussions regarding the viability of the program. Here we will discuss our process and how such an IoT system would apply to the real world. We will use our data and some research to reflect the usability of an IoT system where self-driving cars can communicate with each other.

4.2.1 Self evaluation

As mentioned earlier, we used inspiration from two agile frameworks: scrum and kanban, but chose to lean more towards scrum in the end. We felt the use of scrum helped us reach our goals. However, we could have included our external supervisors more in the scrum process in hindsight. We could have done this by including them in sprint retrospective meetings and discussing what our following sprint goals should have been together.

The IoT system can handle more vehicles, but the road model does not fully support functionalities for having more roads than we currently have in our demonstration. Because of the time constraint, we also found it challenging to focus on scalability. We focused on getting a working demonstration rather than making it scaleable.

One challenge was that our group could not meet as much as we had wished because of work. If we had worked more throughout the process, the need for

extra work, in the end, could have been prevented. In the MosCoW method (Figure 2.3), we were able to finish all the requirements in "must have" and "should have" sections, but none of the "can have" sections. All in all, the group was satisfied with the process.

4.2.2 Educational Value

Developing this solution has been challenging, and our group has learned a lot. We all feel that we have evolved into better developers and that we now would be better suited for solving similar projects in the future.

The project description we started with was very open to interpretation, which gave us much room for exploration. We chose to go with an IoT solution using the car Accenture already owned. Ultimately, we also decided to build a new one. This choice has given us significant insight into making IoT systems and ways to set up communication between them. We believe this resulted in a much more exciting demonstration for Accenture to showcase.

Furthermore, combining such a system with artificial intelligence has been an exciting learning possibility. Due to this project's time constraints, we, for instance, did not integrate machine learning with our server. Integrating machine learning on the server could expand its different capabilities and make the project more scalable and applicable to a broader range of scenarios. Such integration is deemed valuable for the further extension of this project.

While we still were in the first development phase [3.1 Phase 1 - Planning and research phase](#), we decided to adopt the code of the previous group instead of developing our own AI model. Consequently, we believed that this was the right choice, as we believe it would have consumed too much time to start from scratch. Although it was a challenge to understand the program the previous group had written, we assumed that this way, we would get to focus more on our demonstration. The importance of documentation has also been emphasized, which helped our group to write our report.

Our group had little to no experience working with an agile work methodology. In retrospect, we believe we benefitted from this choice. Our daily stand-ups allowed us to have a clear vision of the group's collective challenges. Although we did not implement all aspects of scrum-development or kanban-development, we have gained insight into how a small development team can structure and plan out its workflow in an agile manner.

4.2.3 Real world application

The best way to implement self-driving cars in society is a topic that will take a long time to explore fully. Through our work with this proof-of-concept, we believe we have made an addition to this discussion. Due to time constraints, we have chosen not to create a system that will be viable for all scenarios. Such a system might still seem out of reach. However, individual contributions and new additions to this discussion might take us one step closer to a complete system.

4.2.4 Edge computing and AI

When self-driving vehicles become more prominent and the 5G network becomes more available, there could be a possibility for IoT systems to handle traffic management. Therefore, our group researched how the system will extend to the real world's applications.

The IoT systems usually follow the fog or edge computing architecture with distributed or even decentralized concepts to prevent overloading on servers handling vast loads of data (**iot_platforms**).

One solution to data distribution is that each road has one server responsible for its respective road. If a road is long, it will split into geographical areas where one server has responsibility for their geographical area. Intersections will have a server handling information from both roads' respective servers since the information from both servers is needed to make decisions.

In traffic, there are a lot of unforeseen situations that can happen. Traditional coding will not be able to cover every outcome in a traffic situation. Therefore there will be a need for AI on the servers and the cars. The AI will probably need to train in a safe test environment just like the autonomous vehicles before they are ready for existing road systems.

In our IoT program, the server gives commands to the cars, which overrides the cars' AI. In contrast, ideally, an interaction between the AI on the server and the cars is required for real-world scenarios.

4.2.5 System security

Security and privacy are important topics for any IoT system. Because these systems gather and work with vast amounts of data, they are naturally prone to be attacked. Moreover, as the systems grow and become more interconnected, with many devices worldwide, the imposed risk of such an attack increases drastically (**iot_risk**). Anonymization of personal data, securing connections, and an intruder detection system are all security measures that require attention. The cars must be able to drive both with or without the system, which is an essential feature of system security. If there is a need for the server to shut down, the cars would need to be able to drive independently of the system.

PART III

PRODUCT DOCUMENT

The product documentation is a technical description of the product. The reader of this documentation is assumed to have a technical background. The resulting product is a semi-physical demonstration, and hence, no user interface has been developed. This part serves as a detailed explanation of the structure of the program.

Chapter 5

Program description

The program is an IoT system where vehicles can communicate with a server. The server performs calculations to decide the behavior of the vehicles. The program's primary goal is to increase traffic flow and prevent traffic congestion. The system developed in this project has focused on incoming traffic at an intersection. [Figure 5.1](#) shows a flow diagram that describes the server in our demonstration, mentioned in [3.4.3 Construction of semi physical demonstration](#).

Our solution adheres to the requirements Accenture set for us. As per discussed in [1.2.3 Goals and requirements](#) the vehicles can also drive independently. Moreover, the server can establish multiple connections and determine the optimal speed of each vehicle depending on the incoming situation at an intersection. During the demonstration shown in [3.4.3 Construction of semi physical demonstration](#), it became apparent that the outcome differs; when the cars were driving solely on the integrated AI in contrast to being connected to the server.

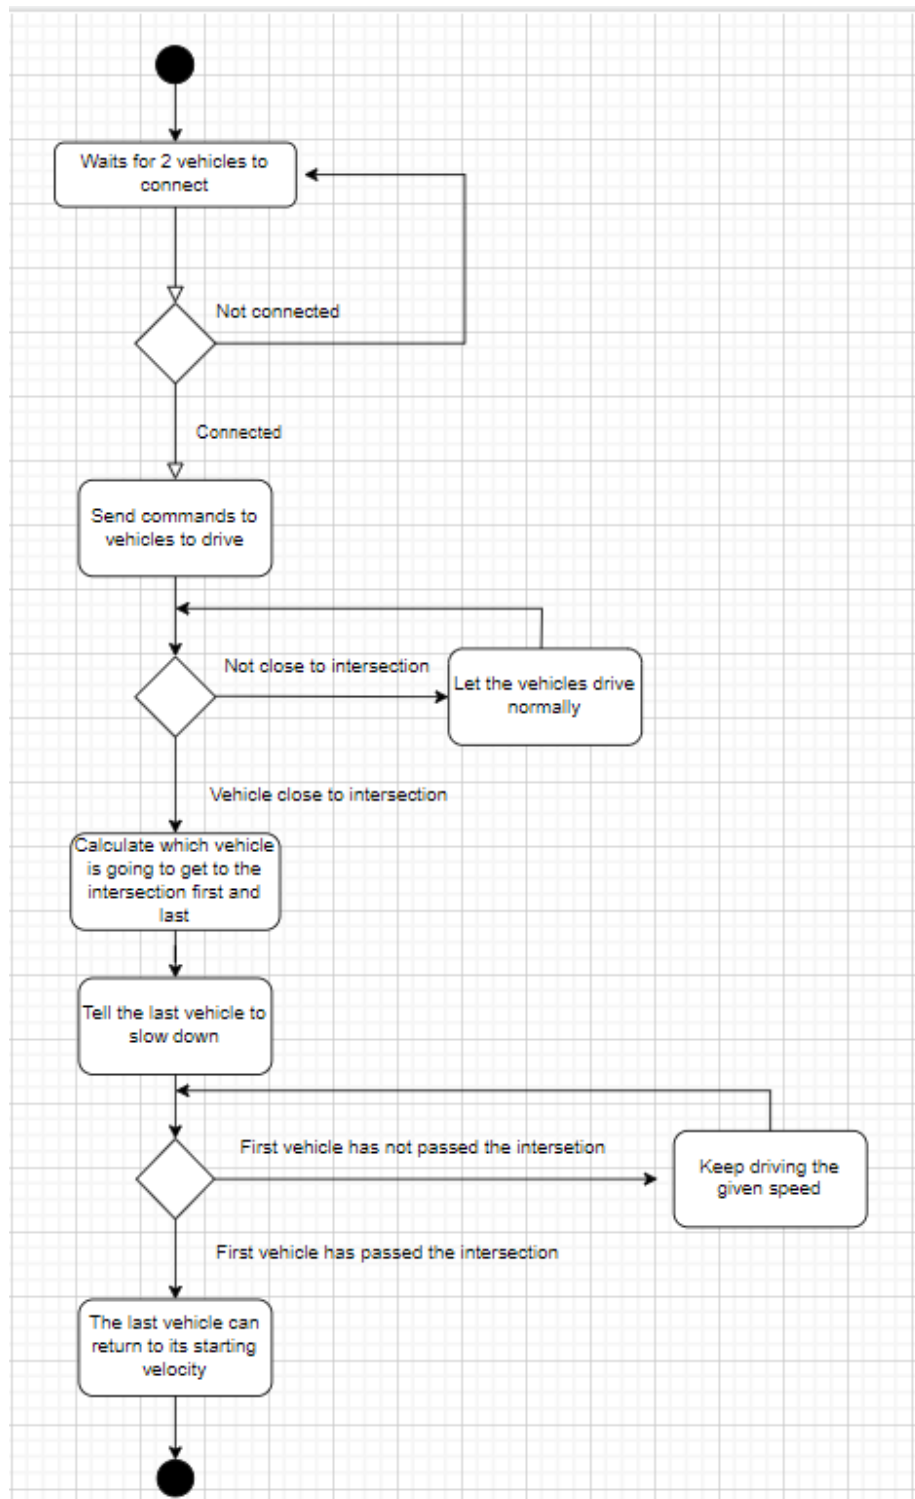


Figure 5.1: Flow diagram for the server. This diagram specifically shows the flow of the semi physical demonstration.

5.1 Essential code snippets behind the system

The central part of this project is to develop a client-server communication system to produce a physical simulation of how a centralized system can contribute to an improved traffic flow. Due to the nature of this project, working with physical models, no graphical user interfaces have been developed. Hence, it is necessary to present critical parts of the code that are responsible for such a system to work. This section will elaborate, in detail, on how essential code snippets interact with each to produce the result.

Furthermore, the code written in this project is in the languages Python and C# using Pycharm and Rider IDE, respectively. Therefore, syntax highlighting is used throughout this section to best simulate the same syntax highlighting used in both IDE, respectively.

In addition, some artistic freedom has been used to present the code snippets; the symbol ... has been used to indicate irrelevant code to the current discussion, and the symbol \hookrightarrow simply means that the line of code following \hookrightarrow is on the same line above but is broken up due to lack of space. Also, each code snippet starts with the class and method that belongs to it. The user manual, in [A User manual](#), is written following the same style as this section.

5.1.1 Initialization

Client.py

The client package is an essential module in the Raspberry Pi vehicles for the success of this project. The client class is mainly responsible for connecting, handling, and sending data to the server. The client class does not aim to be used alone but rather as a superclass for other IoT devices. Hence, it should be inherited and handle everything that pertains to client-server communication in the background.

The client's init method does several things: It reads from the config.json file to store the defined host and port it will use to establish a connection to the server.

```
class Client:
    def __init__(self, properties=None, **kwargs):
        ...
        with open("client/config.json") as f:
            config = json.load(f).get('client')
            if config is not None:
                self.__uri = f"://{config.get('host')}:{config.get('port')}
                     $\hookrightarrow$  }/{self.__class__.__name__.lower()}sHub"
                self.__delay = config.get('delay')
        ...
```

Then, it starts a negotiation process with the server, where it receives a connection id that the client will use during its connection lifetime to the hub.

```
class Client:
    def __init__(self, properties=None, **kwargs):
```

```

...
urllib3.disable_warnings()
response = requests.post(f"http{self.__uri}/negotiate?
    ↪ negotiateVersion=0", verify=False)
self.connection_id = response.json().get("connectionId")
self.websocket_uri = f"ws{self.__uri}?id={self.connection_id}"
...

```

The client also gives itself a random id stored as one of its properties. The client's id is also stored on the server and is mainly used to retrieve and update the client's information.

Furthermore, `Client.__init__` also stores a dictionary of events.

```

class Client:
    def __init__(self, properties=None, **kwargs):
        ...
        self.subscribed_events = {
            "disconnect": self.disconnect,
            "force_patch": self.force_patch,
            "continuously_patch": self.continuously_force_patch
        }
        ...

```

Values in this dictionary are references to functions in this class and is used to invoke certain behaviours by the server. For instance, `await Clients.Client(Context.ConnectionId).SendAsync("disconnect");` from the server will call `def disconnect(self)` in the client.

Vehicle.py

The vehicle class contains all the data and methods of the vehicle. Furthermore, `class Vehicle(Client)` inherits the client class which enables `Vehicle` to perform all the necessary operations to establish connection upon initiation. An important remark is that `Client` performs the negotiation to the server using endpoint

`{self.__class__.__name__.lower()}sHub/negotiate?negotiateVersion=0` meaning that through inheritance and initialization of `Vehicle`, the subclass negotiates with the endpoint `vehiclesHub/negotiate?negotiateVersion=0`, which is mapped in `Program.cs` with

```

...
app.MapHub<VehiclesHub>("/vehiclesHub");
...

```

Furthermore, `Vehicle` also reads from `config.json` to define its initial properties with the snippet shown below:

```

class Vehicle(Client):
    def __init__(self, properties=None, **kwargs):
        ...
        if properties is None and len(kwargs) == 0:

```

```

        with open("client/config.json") as f:
            config = json.load(f).get('vehicle')
            if config is not None:
                self.properties.update(config)
    ...

```

Vehicle also utilizes the property builder of Client.

```

class Vehicle(Client):
    def __init__(self, properties=None, **kwargs):
        ...
        self.property_builder(
            required={'length', 'height', 'width', 'mass'},
            optional={'velocity': 0, 'position': 0, 'travel_plan': None
                    ↪ },
        )
    ...

```

In short, the property builder is used to define the required properties of the vehicle class. The meaning is to restrict somewhat what data the vehicle class should contain. If one should directly initialize Vehicle without using config.json, one must assign values to length, height, width, and mass.

Lastly, Vehicle adds subscribed events that the server can invoke:

```

class Vehicle(Client):
    def __init__(self, properties=None, **kwargs):
        ...
        self.subscribed_events.update({
            "adjust_velocity": self.adjust_velocity
        })

```

Likewise, as in Client, should other events be required for a vehicle, one can add it to the dictionary as proposed above.

5.1.2 Handshake and listener

After initializing the vehicle class as a client with

```

async def main():
    ...
    client = Vehicle()
    ...

```

then the client's listen method can be called:

```

async def main():
    ...
    listener = asyncio.create_task(client.listen())
    ...

```

The listener method is responsible for handling responses and requests from the server. Hence, it must run concurrently as the vehicle continuously sends data to the server.

When the listener is called, the client performs the following code:

```
class Client:
...
    async def listen(self):
    async with websockets.connect(self.websocket_uri) as websocket:
        self.__websocket = websocket
        await self.__handshake()
        await self.__listen()
...
```

The method first opens a WebSocket connection using the stored URI and stores this as a private variable. Then, a handshake with the server is performed:

```
class Client:
...
    async def __handshake(self, protocol: str = "json", version:
        ↪ int = 1):
        data = self.signalr_encode_message({"protocol": protocol, "
            ↪ version": version})
        await self.__websocket.send(data)
        response = self.signalr_decode_message(await self.__websocket.
            ↪ recv())
        if "error" in response:
            print(response)
        else:
            await self.send_non_blocking("AddClient", self.properties)
...
```

The code above describes the handshake process between the client and the server. First, the client informs the server of the protocols it will use throughout its lifetime. Then, it stores the client, in this case, the vehicle, to the server using the defined properties.

Further elaboration, the client invokes the method

```
public partial class VehiclesHub : Hub
{
...
    public async Task AddClient(JsonDocument jsonDocument) {...}
...
}
```

on the server. This method first creates a vehicle with all the provided information sent by the client

```
public partial class VehiclesHub : Hub
{
```

```

...
public async Task AddClient(JsonDocument jsonDocument)
{
    ...
    var vehicle = Vehicle.Create(jsonDocument);
    ...
}
...
}

```

using the static method defined by the Vehicle model. In addition, it assigns the travel plan to the vehicle by using values defined by config.json from the client:

```

{
    ...
    "vehicle": {
        ...
        "travel_plan": {
            "start": {
                "road": 0,
                "lane_reversed": false
            },
            "end": {
                "road": 0,
                "lane_reversed": false
            }
        }
    }
}

```

Furthermore, the vehicle's current lane is also assigned to track which lane the vehicle is driving on. Then, public class VehiclesHubDatabase adds the Vehicle, together with its connection id Context.ConnectionId for easy retrieval.

Moreover, for the sake of the demo, the server is also instructed to wait for a second vehicle to connect before allowing the vehicles to drive

```

public partial class VehiclesHub : Hub
{
    ...
    public async Task AddClient(JsonDocument jsonDocument)
    {
        ...
        vehicle.Velocity = 0;
        _database.Update(vehicle);
        await Clients.Client(Context.ConnectionId).SendAsync("
            ↪ adjust_velocity", vehicle);
        await WaitForVehicles(vehicle, _database.SpeedLimit, 2);
    }
}

```



```
...
}
```

by first setting the vehicle's velocity to zero, updating the new velocity in the database, and adjusting the velocity of the client to zero. Lastly, it calls the `WaitForVehicles` method, which will adjust every client's velocity to the defined `SpeedLimit` in `VehiclesHubDatabase`.

5.1.3 Patch

After initialization and handshake elaborated in [5.1.1 Initialization](#) and [5.1.2 Handshake and listener](#) respectively, the Raspberry Pi vehicles start to drive into an intersection simultaneously. Throughout the journey, the cars are continuously patching to the server by calling the client's `async def send_patch` method.

```
class Client:
    ...
    async def send_patch(self, **kwargs) -> None:
        if self.properties_has_changed(**kwargs) or self.
            ↪ __continuously_patch:
            await self.send_invocation("patch", self.properties)
        else:
            await asyncio.sleep(self.__delay)
```

As seen above `send_patch` calls the `async def send_invocation` method, which communicates the vehicle's current information by invoking `public async Task Patch` on `VehiclesHub`.

The patch method on `VehiclesHub` is responsible for handling the behavior, explicitly adjusting the velocity of individual vehicles:

```
public partial class VehiclesHub : Hub
{
    ...
    public async Task Patch(JsonDocument jsonDocument)
    {
        var vehicle = Vehicle.Create(jsonDocument);
        _database.Update(vehicle);
        vehicle = _database.Fetch(vehicle);
        ...
    }
}
```

The snippet above shows that the method first creates a new vehicle using the information provided by the `Client`. However, since this new vehicle does not contain all the information, such as the travel plan, the method first updates the existing vehicle in the database to refresh the vehicle with the available information. It then fetches the exact vehicle that was stored in the handshake, mentioned in [5.1.2 Handshake and listener](#). When the vehicle is successfully retrieved, it will then handle this vehicle accordingly:

```

public partial class VehiclesHub : Hub
{
    ...
    public async Task Patch(JsonDocument jsonDocument)
    {
        ...
        await HandleIntersection(vehicle);
        await HandleInsideIntersection(vehicle);
        await HandleEndOfRoute(vehicle);
        ...
    }
}

```

Shortly summarized `HandleIntersection` is responsible for adjusting the velocity of every vehicle approaching the intersection to avoid collisions. Furthermore, `HandleInsideIntersection` increases the speed to `VehiclesHubDatabase` defined `SpeedLimit`. Lastly, `HandleEndOfRoute` ensures that any vehicle that has completed its journey, defined during the handshake, terminates its connection with the server.

5.1.4 VehiclesHubDatabase

The `VehiclesHubDatabase` played a vital role in this project. As discussed on [3.3 Phase 3 - WebSocket with SignalR](#), the project needed a database that could handle the continuous changes in each `Vehicle` position in real-time. During the research, the group could not conclude any databases that would fit our requirements. Hence, `VehiclesHubDatabase` was created to serve as a live in-memory database that could continuously update the positions of each vehicle on each lane.

Before starting with `VehiclesHubDatabase`, it was required to define what a road, lane and intersection is, respectively. Thus, `Road.cs`, `Lane.cs` and `Intersection.cs` was developed. Consequently, `VehiclesHubDatabase` was created.

```

public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    private readonly Intersection _intersection;

    private readonly HashSet<Vehicle> _vehicles = new();
    private readonly HashSet<Lane> _lanes = new();
    public int Count => _vehicles.Count;
    private readonly Dictionary<string, Vehicle> _connectionIds =
        => new();
    private readonly Dictionary<Vehicle, string>
        => _vehiclesConnectionId = new();
    ...
    public double SpeedLimit => 80;
    ...
}

```




Figure 5.2: This figure shows two roads each with two lanes and a total length of 300cm. The overlapping part forms a square which is the intersection. This configuration was heavily considered when creating representational models on the server, and also used for the demo as a result.

Currently, `VehiclesHubDatabase` only holds one intersection, due to time constraint this was not extended for a configuration with multiple intersections. Moreover, both vehicles and lanes are stored inside a hashset for fast retrieval. In addition, `Count` is used in `WaitForVehicles`, elaborated in sub-section 5.1.2 [Handshake and listener](#), and `_connectionId` and `_vehiclesConnectionId` is used during `Patch` to invoke `adjust_velocity` on individual vehicles. While, `SpeedLimit` defines the upper speed vehicles are limited to on the two roads shown in [Figure 5.2](#).

The road configuration found in [Figure 5.2](#) is defined in the constructor using a builder pattern.

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public VehiclesHubDatabase()
    {
        _intersection = new Intersection().
```

```
        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane()).
        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane());

        _intersection.ConnectedLanes().
            ForEach(lane => _lanes.Add(lane));
    }
    ...
}
```

Lastly, `VehiclesHubDatabase` is added as a singleton service in `Program.cs` to ensure that we have a static database throughout the lifetime of the program:

```
...
builder.Services.AddSingleton<IVehiclesHubDatabase>(new
    ↪ VehiclesHubDatabase());
...
```

It is also worth to mention some of the core functionalities of `VehiclesHubDatabase`:

Adding vehicles

Calling the `Add` method makes it possible to add vehicles:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public void Add(Vehicle vehicle, string? connectionId = null)
        ↪ {...}

    public void Add(Vehicle vehicle, Lane? lane = null) {...}
    ...
}
```

Removing vehicles

Removing vehicles can be achieved by calling the `Remove` method:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public void Remove(Vehicle vehicle) {...}
    ...
}
```

Updating vehicles

Updating either a specific information of a vehicle or all vehicles can be done by calling the `Update` method:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public void Update(Vehicle? vehicle = null) {...}
    ...
}
```

Fetching vehicles

One can fetch an existing vehicle from the database by passing a vehicle with the same GUID with:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public Vehicle? Fetch(Vehicle vehicle) {...}
    ...
}
```

Get the connection Id of a particular vehicle

By passing a vehicle into

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public string? ConnectionId(Vehicle vehicle) {...}
    ...
}
```

one can retrieve the connection Id that the given vehicle is using.

Find vehicles approaching the intersection

Maybe the most important feature of `VehiclesHubDatabase` is the two method shown below:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public IEnumerable<Vehicle> NextVehiclesIn() {...}
    public IEnumerable<Vehicle> OnlyFirstIntoNextVehiclesIn() {...}
}
```

The first method `NextVehiclesIn` returns a list of vehicles currently approaching the intersection defined in the constructor, ordered with respect to the vehicle closest to the intersection. The latter method `OnlyFirstIntoNextVehiclesIn` returns an ordered list of the closest vehicle per lane.

5.1.5 RoutePlanner and SetTravelPlan

`RoutePlanner` is a class that determines and keep track of the `Vehicle`'s journey. `RoutePlanner` saves segments of lanes and intersections as nodes in a linked list:

```
public class RoutePlanner
{
    private LinkedList<IRoadComponent> _visitedNodes = new();
    private LinkedList<IRoadComponent> _travelPlan = new();
    ...
}
```

The variable `_travelPlan` stores all the nodes that `Vehicle` will traverse through in succession. Whenever `Vehicle` has traversed the whole length of a node, the node moves to the tail of `_visitedNodes`. Thus, it provides an easy way to keep track of the `Vehicle`'s whereabouts.

During the handshake, described in 5.1.2 Handshake and listener, `Vehicle` builds the `_travelPlan` through its `SetTravelPlan` method:

```
public class Vehicle : IDevice
{
    ...
    public void SetTravelPlan(Lane startLane, Lane? endLane = null)
    {
        var reversed = startLane.Reversed;
        var startNode = startLane.Node();
        var intersections =
            reversed ?
                startLane.CurrentRoad()?.Intersections.OrderByDescending(
                    ↪ kvp => kvp.Key) :
                startLane.CurrentRoad()?.Intersections.OrderBy(kvp => kvp.
                    ↪ Key);
        ...
    }
    ...
}
```

By taking in the `startLane` and `endLane`, representing where `Vehicle` should start and end its journey, it first segments the `startLane` into a `LaneNode` and then find all the intersections on this current `Lane`, and order them chronologically depending of the lane is `Reversed` or not. `SetTravelPlan` will further on iterate through all the `Intersections` and append a `LaneNode` and an `IntersectionNode` to `RoutePlanner` for each `Intersection`.

```
public class Vehicle : IDevice
{
    ...
    public void SetTravelPlan(Lane startLane, Lane? endLane = null)
    {
        ...
    }
}
```

```

var prevPos = reversed ? startLane.Length : 0.0;
Intersection? intersectionConnectedToEndLane = null;

if (intersections != null)
    foreach (var (position, intersection) in intersections)
    {
        startNode.Length = Math.Abs((int) (prevPos - position - (
            ↪ reversed ? intersection.Length : 0)));
        _route.AddComponent(startNode).AddComponent(intersection.
            ↪ Node());
        prevPos = position + (reversed ? 0 : intersection.Length);
        if (endLane == null) continue;
        if (!intersection.ConnectedRoads.ContainsKey(endLane.
            ↪ CurrentRoad() ?? new Road())) continue;
        intersectionConnectedToEndLane = intersection;
        break;
    }
    ...
}
    ...
}

```

The above loop will iterate until one of the `Intersection` is connected to the `endLane` or until all the `Intersections` are exhausted. In the end, `SetTravelPlan` will append the last segment of `startLane` or the remaining segment of `endLane` should `endLane` either be defined, and connected to one of `startLanes` intersections, or is equal to `startLane`.

The main reason for creating `RoutePlanner` in this project was mainly to answer the following questions:

- Is `Vehicle` currently inside an intersection?
- Is `Vehicle` currently on a lane?
- Which intersection is the next intersection for this `Vehicle`?
- How far is `Vehicle` away from the next intersection?

By using `RoutePlanner`, we were able to answer the questions above with the following implementations respectively:

```

public class Vehicle : IDevice
{
    ...
    public bool InIntersection() =>
        _route.CurrentNode()?.Value is IntersectionNode;
    public bool OnRoad() =>
        _route.CurrentNode()?.Value is RoadNode or LaneNode;
    public IntersectionNode? NextIntersection() =>
        OnRoad() ? _route.NextNode()?.Value as IntersectionNode : null
            ↪ ;
    public double ToNextIntersection()

```

```
{
    if (InIntersection())
        return 0;
    if (OnRoad() && NextIntersection() != null)
        return Math.Abs(Position - _route.DistanceWithCurrentNode);
    return -1;
}
...
}
```


Appendix A

User manual

We have written a manual for people who want to recreate our demonstration. The demonstration could for example be shown off at exhibitions. The manual could also be used for people who want to further test and develop our IoT-system.

First check the IP-address of the internet you are connected to, and the usable ports. Make sure that your computer hosting the server and the vehicles are connected to the same network.

```
$ipconfig getifaddr en0
192.168.56.208
```

Then open the Server solution in your code editor, we have used visual studio. Under the folder “properties” there is a file called launchSettings.json. In that file write in the ip address and the port in the applicationUrl-section:

```
"profiles": {
  "SignalRServer": {
    "commandName": "Project",
    "dotnetRunMessages": true,
    "launchBrowser": false,
    "applicationUrl": "https://192.168.56.208:7058;http://192.168.
      ↪ 56.208:5048",
    "environmentVariables": {
      "ASPNETCORE_ENVIRONMENT": "Development"
    }
  },
}
```

After that, open the Client solution. Here we used Pycharm as the IDE. Open the config.json document and replace the same IP address and port:

```
"client": {
  "host": "192.168.56.208",
  "port": 5048,
  "delay": 0.1
},
```

If the vehicles have not connected to that network before, they need to log on to that network. To log on to a new network; connect the Raspberry Pi to a power source and a display, and use the user interface. The display port on the Raspberry Pi is a micro USB port. When the Raspberry Pi has booted up, click on the internet icon and connect to the same network as the server. If the network has been connected to it before, Raspberry Pi will automatically connect to that network during boot up.

The initial speed of the vehicles can be changed by adjusting the value of `SpeedLimit` located at `VehicleHubDatabase` under the `Database` folder on the server solution:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public double SpeedLimit => 80;
    ...
}
```

Moreover, the configuration of the road can also be changed by changing the following section of the code:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public VehiclesHubDatabase()
    {
        _intersection = new Intersection().
        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane()).
        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane());
        ...
    }
    ...
}
```

The current configuration shown above corresponds to the same configuration shown in [Figure 5.2](#). The intersection can be moved by adding additional arguments such as:

```
public class VehiclesHubDatabase : IVehiclesHubDatabase
{
    ...
    public VehiclesHubDatabase()
    {
        _intersection = new Intersection().
        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane(), 200).
    }
}
```

```

        AddRoad(new Road {Length = 300}.
            AddLane(null, true).
            AddLane(), 150);
        ...
    }
    ...
}

```

Figure A.1 shows the resulting topology by adding extra parameters to the code above.




Figure A.1: An alternative configuration of the intersection and its connected roads by adding extra parameter to the initialization of the intersection as described by the previous code snippet.

The length is in centimeters and needs to correlate with the length of the physical track. We used tape to show where the roads were. However, this is unnecessary. We recommend a track between three to four meters long for the best results.

Lastly, position both vehicles down at the start of the two tracks, turn on the server and connect to the power banks. The vehicles should automatically connect to the server after 20-30 seconds. When both vehicles have established a connection to the server, the demonstration has started.

Appendix B

Glossaries and Acronyms

B.1 Glossaries

Application programming interface - An API, or application programming interface, is a set of rules that define how applications or devices can connect to and communicate with each other (**rest_api**).

Artificial intelligence - A field which combines computer science and robust datasets to enable problem solving (**artificial_intelligence**).

Cloud computing - Delivery of IT resources over the internet such as data storage, servers and software.

Edge computing - A distribution architecture where processing and data storage are brought closer to the data source to prevent servers overloading from handling vast loads of data.

Hypertext transfer protocol - A protocol for transferring data between network devices.

InfluxDatabase - A time series-database which provides SQL-like syntax that is quick to query resources. Time-series database automatically includes timestamp for each database entry.

Internet of Things - The Internet of Things refers to physical objects that communicate using sensors, cameras, software, or other technologies connecting and exchanging data.

Internet of Vehicles - An IoV system is a distributed system for wireless communication and information exchange between vehicles through agreed-upon communication protocols (**chinese_iov**).

Raspberry Pi - A series of small single-board computers developed by the Raspberry Pi foundation.

Remote procedure call - A protocol that provides the high-level communications paradigm used in the operating system. RPC presumes the existence of a low-level transport protocol, such as Transmission Control Protocol/Internet

Protocol (TCP/IP) or User Datagram Protocol (UDP), for carrying the message data between communicating programs. (**rpc_ibm**)

Representation state transfer - An architectural style when used together with API, provides communication through HTTP.

Tensor Processing Unit - Tensor Processing Units (TPUs) are Google's custom-developed application-specific integrated circuits (ASICs) used to accelerate machine learning workloads (**google_tpu**).

Transmission control protocol - A protocol that ensures reliable delivery of data over the internet.

Websockets - Websockets is a protocol that provides a bidirectional communication between clients and server by establishing a single TCP connection in both direction (**rfc_websockets**).

B.2 Acronyms

AI - Artificial intelligence

API - Application programming interface

CRUD - Create, read, update, and delete

HTTP - Hypertext transfer protocol

IDE - Integrated Development Environment

IoT - Internet of things

IoV - Internet of vehicles

REST - Representational state transfer

RPC - Remote procedure call

SQL - Structured query language

TCP - Transmission control protocol

TPU - Tensor processing unit

Appendix C

Sprint documents

28.03.2022

Retrospective: (For sprint 3)

Hva gikk bra?

- Fikk til å koble bilene til server
- Hele gruppen har større forståelse av koden
- SignalR-serveren fungerer bra
- Hatt bedre kommunikasjon med veiledere, spesielt med intern

Hva kunne gått bedre?

- Skjevfordeling av tekniske oppgaver
- Blitt dårligere på å be om teknisk hjelp av eksterne veiledere

Hvilke tiltak kan vi gjøre?

- Fordele oppgaver annerledes fremover.
- Fortsette å sende meldinger til veiledere, spesielt når vi sitter fast med noe teknisk.

Figure C.1: A snapshot of a sprint retrospect perspective meeting from our log. There are three sections: What went well? What could have gone better? And what can we do for next time?

Sprint planning for sprint 4 (21.03 - 04.06)

- Oppgaver fra backlog
 - **Disposisjon til bacheloroppgave**
 - **Teori til bacheloroppgave**
 - **Få server til å sende kommando til flere biler.**

Mulige oppgaver:

- **Enhetstesting**
- **Simulering**

Figure C.2: A snapshot of a sprint planning meeting from our log. The upper section contains the tasks we are going to do from the backlog. The section under contains possible but not necessary tasks.